

Carlos Urteaga Reyesvera

AI Delivery Lead – AI Architect

Mexico City, Mexico • carlos.urteaga@outlook.com • LinkedIn • GitHub • urteaga.tech

Professional Summary

AI Delivery Lead and AI Architect with 10+ years of experience delivering production-oriented AI systems across consulting, financial services, telecommunications, retail, education, and technology. Works across architecture, delivery, and hands-on engineering to translate ambiguous business requirements into enterprise RAG systems, AI agents, document intelligence workflows, and MLOps / LLMOps capabilities. Recent work spans Azure and AWS, AI evaluation and observability, infrastructure automation, release discipline, and cross-functional execution across engineering, QA, DevOps, data science, and client teams.

Education

Ph.D. Student in Computer Science, Universidad Americana de Europa (UNADE) Remote, expected 2027
Research focus: maturity, evaluation, governance, AgentOps, and operational documentation of LLM-powered AI agents.

M.Sc. in Computer Science, Instituto Tecnológico Autónomo de México (ITAM) Mexico City, Mexico
Degree completed; professional license issued.

B.S. in Telematics Engineering, Instituto Tecnológico Autónomo de México (ITAM) Mexico City, Mexico
Degree completed; professional license issued.

Experience

Globant

Hybrid, Mexico City

Mar 2019 – Present

Data Scientist III & AI Focal

Jan 2023 – Present

Internal progression: Data Scientist I (Mar 2019 – Apr 2021) → Data Scientist II (Apr 2021 – Dec 2022). Operated in practice as Tech Lead / AI Delivery Lead across multiple initiatives since 2021.

Azure document intelligence project (NDA)

- Led delivery for enterprise AI initiatives across Azure and AWS, coordinating a cross-functional team of approximately 10 people across engineering, QA, DevOps, data science, architecture, and client developers.
- Designed a document-intelligence agent for audit and risk workflows that reduced processing time from 500+ hours to approximately 30 minutes using Azure Document Intelligence and SME-guided prompts.
- Implemented workflows that classify 8+ document types, generate audit findings, and extract risk information through structured prompts and document-processing pipelines.
- Introduced staging, QA gates, controlled promotion, and release discipline to improve delivery readiness in a critical client environment.
- Built and validated Azure infrastructure automation for LLMs, agents, private endpoints, multi-region deployments, and environment-specific configurations.
- Established an agent-assisted SDLC using low-level design artifacts, contextual Markdown guidance, and operational prompts for coding agents, increasing squad throughput by at least 3x in user stories while reducing avoidable refactoring.
- Owned prioritization of user stories and refinement of delivery scope based on client needs, aligning execution with business and operational objectives.
- Stabilized delivery capacity through technical coaching and developer support during a critical phase, helping retain a key team member at risk of leaving the project.

G4G Fast Ops

- Led development of an AWS-based Slack agent with online and batch workflows, Dagster orchestration, Amazon SQS, CloudWatch, action traceability, delegated identity, and initial evaluation with RAGAS.
- Recovered a delayed AI project between September and October 2025 by defining a structured remediation path across prompts, data, and model configuration, restoring expected model behavior and enabling delivery.

Enterprise RAG and AI platform work

- Rebuilt an enterprise chatbot by redesigning its RAG architecture and orchestration, improving quality from 50% to 90% without replacing the underlying model.
- Designed and deployed RAG services with Azure, AWS, Python, LangChain, and FastAPI for secure retrieval across enterprise documents, presentations, and spreadsheets.
- Fine-tuned a GPT-2 model for semantic classification of financial statements, achieving a 20% performance improvement.
- Validated enterprise readiness for AI agents through load, scalability, and resiliency testing at 3,000, 5,000, 10,000, and up to 30,000 concurrent users.
- Supported technical pre-sales, AI capability assessment, engineering standards, and talent evaluation to align delivery practices with broader organizational needs.

Data Scientist II

Apr 2021 – Dec 2022

- Led adoption of MLOps practices for internal ML systems, introducing CI/CD/CT, continuous training, monitoring, and release discipline to improve reliability, observability, and reproducibility.
- Built a Random Forest model to predict candidate seniority with 90% accuracy and led a multidisciplinary team using Kanban to deliver classification workflows aligned with hiring operations.
- Converted notebook-based ML solutions into production REST APIs, improving deployment readiness and completing one delivery one month ahead of schedule.
- Designed Python, Flask, Celery, and Docker workflows for prediction services and streaming data processing.
- Defined an MLOps strategy for talent-acquisition and evaluation models, using MLflow and Argo Workflows to support monitoring, maintenance, and optimization over a two-year period.
- Led the transition from PoC to MVP in two weeks for an interviewer-recommendation platform, defining backend and ML operating flow for production adoption.
- Integrated Google Calendar, time zones, languages, and availability constraints to improve interviewer matching and enable round-robin allocation across gatekeepers.
- Reduced interviewer assignment time from approximately 2 days to 10 minutes and expanded gatekeeper availability from 2 to 40 per technology.
- Implemented early lifecycle practices including canary-style rollout patterns, queue management, and asynchronous processing with Celery.
- The solution reached production use inside Globant, supported approximately 1,000 candidates or interviews per month, and served a network of around 8,000 gatekeepers across an organization present in 33 countries.

Data Scientist I

Mar 2019 – Apr 2021

- Built a real-time face-recognition and people-counting solution with Python, TensorFlow, Keras, and Flask, achieving approximately 90% operational accuracy at around 20 FPS.
- Led telecom cost analysis for a U.S. population of approximately 100,000 users, identifying savings opportunities of up to 25% per line and approximately 30% in plan-reassignment cases based on usage, calling patterns, and location.
- Designed analytical outputs that helped frame mobile-plan optimization as a measurable cost-efficiency opportunity.
- Developed an XGBoost-based domestic-violence recurrence model with Python, Docker, NLP, and feature engineering for initiatives in Uruguay, Brazil, and Honduras, achieving 82% accuracy.
- Designed ETL pipelines and Flask APIs to migrate and standardize retail catalog data into Snowflake, consolidating heterogeneous product information from approximately 250,000 merchants across Latin America, including implementations for Argentina and Colombia.
- Built extraction and normalization pipelines with Pandas, scikit-learn, regular expressions, and Docker, improving extraction quality by 80% when source data was incomplete.

Compartamos Banco / Gentera

Head of Data Science

Sep 2018 – Mar 2019

Mexico City

- Led the architecture and delivery of machine-learning products for financial services using Python, Spark, Hadoop, H2O, and SPSS Modeler.
- Managed a team of approximately five people with responsibility for hiring and prioritization of business needs.
- Moved local processing workloads into a hybrid AWS architecture to support big-data requirements.
- Automated training, feature-selection, and model-tuning pipelines following CRISP-DM practices.
- Migrated up-selling, cross-selling, churn, and credit-risk models to production, reducing execution time from 15 days to 10 hours.
- Improved behavioral models by 0.15 ROC-AUC points and increased a churn model's accuracy from 82% to 95%.
- Designed and productionized a credit-risk model on Spark and Hadoop, covering feature engineering, training, tuning, validation, and deployment.
- Designed an A/B-style campaign-validation approach using comparable branches and states, historical similarity, and population density criteria to support a 10% user-growth objective worth approximately MXN 1.5M per group over 3 months.
- Improved a bulk extraction-and-load process into the production database, reducing runtime from approximately 1.5 days to 40 minutes.
- Quantified that client-retention improvements represented approximately 15% of credit placement value, helping prioritize loyalty and campaign initiatives.

Telcel

May 2018 – Aug 2018

Senior Data Scientist

Mexico City

- Analyzed and cleaned IVR data for approximately 82,000 users using Python, Hadoop, and SQL.
- Redesigned an IVR mobile experience to reduce incoming calls and contact-center costs.
- Built Tableau dashboards for retention, up-selling, cross-selling, churn, and marketing analysis.
- Developed tracking visibility that distinguished real retention from an initial 10% estimate, adjusting it to 2%.

KPMG

Oct 2017 – Apr 2018

Senior Information Security Analyst

Mexico City

- Applied machine learning and Python to internal information-security risk analysis, improving detection by approximately 5%.
- Developed Random Forest analysis at national scale to verify IT policy compliance and detect process leakage.
- Automated incident analysis and reporting, reducing a process from 15 days to approximately 2 hours.

Earlier Research, Teaching, and Engineering

Instituto Tecnológico Autónomo de México (ITAM)

Aug 2010 – May 2018

Lecturer, Teaching Assistant, and Laboratory Assistant

Mexico City

- Taught and supported courses in computer networks, digital systems, mechatronics, robotics, and introductory engineering.
- Guided projects using Python, C, ROS, Arduino, Quartus, computer vision, and autonomous robots.
- Administered seven teaching and research laboratories across Linux, macOS, and Windows environments.

ITAM and Microsoft Research

Jan 2016 – Dec 2016

Data Scientist / Research Assistant / Project Manager

Mexico City

- Consolidated weekly activity logs and screenshots for a study on multitasking and academic performance.
- Automated a process equivalent to approximately one month of manual work using Python.
- Applied computer vision, machine learning, and neural networks to classify productive vs. distracting activities.

Instituto Federal de Telecomunicaciones (IFT)

Dec 2014 – Jun 2015

Senior Analyst / Senior iOS Developer

Mexico City

- Estimated internet service demand in small localities and proposed GIS-based approaches for fiber-route design.
- Built an iOS app in Objective-C to capture antenna identifiers, signal strength, and mobile-network measurements for service-quality auditing.

Certifications

- Microsoft Certified: Azure AI Engineer Associate (Microsoft, 2026)
- Claude Code in Action (Anthropic, 2026)
- Model Context Protocol (Anthropic, 2026)
- AI Agents in LangGraph (DeepLearning.AI, 2025)
- AI Agents Fundamentals (Hugging Face, 2025)
- Machine Learning Engineering for Production (MLOps) Specialization (DeepLearning.AI, 2022)

Selected Publications

- Urteaga-Reyesvera, J.C., Lopez Murphy, J.J. (2026). *Agent Cards: A Documentation Standard for Operational AI Agents. Advances in Computational Intelligence. MICAI 2025 International Workshops*, LNCS 16265, Springer.
- Urteaga-Reyesvera, J.C., Cadena Martínez, R. (2026). *Ask First, Test on Demand: A Deference-Gated Socratic Agent Design. Advances in Computational Intelligence. MICAI 2025 International Workshops*.
- Urteaga-Reyesvera, J.C., Possani-Espinosa, A. (2016). *Scorpions: Classification of Poisonous Species Using Shape Features*. 2016 International Conference on Electronics, Communications and Computers (CONIELECOMP), IEEE.

Skills

Leadership & delivery: enterprise AI delivery, technical leadership, cross-functional execution, stakeholder alignment, production readiness, prioritization, technical pre-sales.

Generative AI: RAG, AI agents, document intelligence, LangChain, LangGraph, prompt engineering, RAGAS, LLM evaluation and observability, LLMOps, AgentOps.

Engineering & platform: Python, FastAPI, Flask, REST API, CI/CD/CT, MLflow, Dagster, Docker, Celery, monitoring, load testing, infrastructure automation, IaC.

Cloud & data: Azure AI, Azure AI Search, Azure Document Intelligence, Azure Functions, AWS Lambda, Amazon SQS, Amazon CloudWatch, Snowflake, GCP, ETL, feature engineering, NLP, computer vision, XGBoost, TensorFlow, Keras.